

Day 1. Settings

NPEX Reinforcement Learning

2020.08.27.

Jaeuk Shin



CORE
Control + Optimization Research Lab

Contents

- Package Installation

- Anaconda
- PyTorch
- OpenAI Gym

- Examples

- PyTorch
- OpenAI Gym

Anaconda



www.anaconda.com/

Handy management of separate environments

Convenient installation of scientific computing libraries →
numpy, scipy, pandas, matplotlib, tensorflow, PyTorch, ... etc.

Anaconda

Create your environment:

```
sju5379@sju5379-System-Product-Name:~$ conda create -n npex python=3.6
Collecting package metadata (current_repodata.json): done
Solving environment: done
```

Check whether the environment is successfully created:

```
sju5379@sju5379-System-Product-Name:~$ conda env list
# conda environments:
#
base                *  /home/sju5379/anaconda3
baselines           /home/sju5379/anaconda3/envs/baselines
drl                 /home/sju5379/anaconda3/envs/drl
npex                /home/sju5379/anaconda3/envs/npex
torchbeast          /home/sju5379/anaconda3/envs/torchbeast
tsearch             /home/sju5379/anaconda3/envs/tsearch
```



Anaconda

Activate the created environment:


```
sju5379@sju5379-System-Product-Name:~$ conda activate npex  
(npex) sju5379@sju5379-System-Product-Name:~$ conda install
```

```
(npex) sju5379@sju5379-System-Product-Name:~$ conda list  
# packages in environment at /home/sju5379/anaconda3/envs/npex:  
#  
# Name                        Version           Build    Channel  
_libgcc_mutex                0.1              main  
blas                         1.0              mkl  
ca-certificates              2020.6.24        0  
certifi                      2020.6.20        py36_0  
cloudpickle                  1.3.0            pypi_0   pypi
```

For more info:

<https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>

PyTorch



High-level Deep Learning Framework

Efficient Building/Training of large-scale models (ex. OpenAI GPT-3)

No such DL frameworks? → multiprocessing, CUDA, etc.

PyTorch - Installation

In your Conda env:

conda install pytorch torchvision cpuonly –c pytorch

```
(npex) sju5379@sju5379-System-Product-Name:~$ conda install pytorch torchvision  
cpuonly -c pytorch
```

For the class, we will use cpu-only version(if you have already installed gpu version, it doesn't matter)

pytorch.org/



CORE
Control + Optimization Research Lab

PyTorch - Examples

Open `torch_test.py`







When computing the forwards pass, autograd simultaneously performs the requested computations and builds up a graph representing the function that computes the gradient.

Note : graph is recreated from scratch at **every iteration!**

[Reference] pytorch.org/docs/stable/notes/autograd.html

OpenAI Gym

provides various types of MDPs/RL benchmarks



OpenAI Gym

high-level API for agent-environment interaction

Intuitive abstractions, easy usage

ex. MuJoCo API vs Gym API

- ← MUJOCO.ORG
- Overview
- Computation
- Modeling
- Programming
- XML Reference
- API Reference
- HAPTIX
- Unity Plugin
- Introduction
 - Key features
 - Model instances
 - Examples
- Model elements
 - Options
 - Assets
 - Kinematic tree
 - Stand-alone
- Clarifications
 - Not object-oriented
 - Softness and slip
 - Types, names, ids
 - Bodies, geoms, sites
 - Joint coordinates



Preface

This is an online book about the MuJoCo physics simulator. It contains all the information needed to use MuJoCo effectively. It includes introductory material, technical explanation of the underlying physics model and associated algorithms, specification of MJCF which is MuJoCo's XML modeling format, user guides and reference manuals. Additional information, answers to user questions as well as a collection of models can be found on the [MuJoCo Forum](#).

Chapter 1: Overview



OpenAI Gym - Installation

In your Conda env:

`pip install gym`

```
(npex) sju5379@sju5379-System-Product-Name:~$ pip install gym
Processing ./cache/pip/wheels/be/a1/84/6b4caa6c1cea703acbfea8a24cc3c1729bd359cd
4a65755d8b/gym-0.17.2-py3-none-any.whl
Collecting scipy
  Downloading scipy-1.5.2-cp36-cp36m-manylinux1_x86_64.whl (25.9 MB)
    |████████████████████████████████████████| 25.9 MB 1.3 MB/s
```

gym.openai.com

github.com/openai/gym




CORE
Control + Optimization Research Lab

OpenAI Gym - Examples

Pendulum Swing Up

Goal : keep a frictionless pendulum standing up



OpenAI Gym - Examples

Open gym_test.py

```
1  import gym
2
3
4  env = gym.make('Pendulum-v0')
5
6  state = env.reset()
7  done = False
8
9
10 while not done:
11     state, reward, done, _ = env.step(env.action_space.sample())
12     env.render()
13
14
15 env.close()
```

Class Env(object):

action_space =
observation_space =

def step(self, action):

def reset(self):

def render(self, mode='human'):

def close(self):



OpenAI Gym - Examples

See pendulum.py

```
class PendulumEnv(gym.Env):
    metadata = {
        'render.modes': ['human', 'rgb_array'],
        'video.frames_per_second': 30
    }

    def __init__(self, g=10.0):
        self.max_speed = 8
        self.max_torque = 2.
        self.dt = .05
        self.g = g
        self.m = 1.
        self.l = 1.
        self.viewer = None
```

Why observation space instead of state space?
Notion of **Partially-Observable MDP(POMDP)**

determine the **state space \mathcal{S}**
and the **action space \mathcal{A}** ←


```
high = np.array([1., 1., self.max_speed], dtype=np.float32)
self.action_space = spaces.Box(
    low=-self.max_torque,
    high=self.max_torque, shape=(1,),
    dtype=np.float32
)
self.observation_space = spaces.Box(
    low=-high,
    high=high,
    dtype=np.float32
)

self.seed()
```

OpenAI Gym - Examples



(a) Sample observation



(b) Layout of the 5×5 maze in (a)

state space?

observation space?



OpenAI Gym - Examples

```
def reset(self):
```

```
    high = np.array([np.pi, 1])
```

```
    self.state = self.np_random.uniform(low=-high, high=high)
```

```
    self.last_u = None
```

```
    return self._get_obs()
```

sample an initial state s_0

```
def _get_obs(self):
```

```
    theta, thetadot = self.state
```

```
    return np.array([np.cos(theta), np.sin(theta), thetadot])
```

state : $s = (\theta, \dot{\theta})$

observation : $o = (\cos \theta, \sin \theta, \dot{\theta})$



OpenAI Gym - Examples

```
def step(self, u):
```

most important method of Gym Env class!

```
    th, thdot = self.state # th := theta
```

current state s_t kept by the env

```
    g = self.g
```

```
    m = self.m
```

```
    l = self.l
```

```
    dt = self.dt
```

```
    u = np.clip(u, -self.max_torque, self.max_torque)[0]
```

```
    self.last_u = u # for rendering
```

```
    costs = angle_normalize(th) ** 2 + .1 * thdot ** 2 + .001 * (u ** 2)
```

reward $r(s_t, a_t)$

```
    newthdot = thdot + (-3 * g / (2 * l) * np.sin(th + np.pi) + 3. / (m * l ** 2) * u) * dt
```

```
    newth = th + newthdot * dt
```

```
    newthdot = np.clip(newthdot, -self.max_speed, self.max_speed)
```

sample s_{t+1}

```
    self.state = np.array([newth, newthdot])
```

```
    return self._get_obs(), -costs, False, {}
```

Given the current state s_t and an action a_t , we receive the reward r_t and the next state s_{t+1} from the environment:

$$s_{t+1} \sim p(\cdot | s_t, a_t), \quad r_t = r(s_t, a_t).$$

Thank you!